

CENG443

Heterogeneous Parallel Programming

Introduction

Işıl ÖZ, IZTECH, Fall 2023

06 October 2023



Programming Model

CUDA (Compute Unified Device Architecture) C General purpose programming model for NVIDIA GPUs

Parallel computing platform and application programming interface (API)

Standard C Code

```
void saxpy(int n, float a,
          float *x, float *y)
{
    for (int i = 0; i < n; ++i)
        y[i] = a*x[i] + y[i];
}

int N = 1<<20;

// Perform SAXPY on 1M elements
saxpy(N, 2.0, x, y);
```

C with CUDA extensions

```
__global__
void saxpy(int n, float a,
          float *x, float *y)
{
    int i = blockIdx.x*blockDim.x + threadIdx.x;
    if (i < n) y[i] = a*x[i] + y[i];
}

int N = 1<<20;
cudaMemcpy(x, d_x, N, cudaMemcpyHostToDevice);
cudaMemcpy(y, d_y, N, cudaMemcpyHostToDevice);

// Perform SAXPY on 1M elements
saxpy<<<4096,256>>>(N, 2.0, x, y);

cudaMemcpy(d_y, y, N, cudaMemcpyDeviceToHost);
```

Development Environment

CUDA Toolkit, nvcc compiler

Your PC with an NVIDIA GPU - Preferable

Google Colab

Amazon Cloud (AWS credit from NVIDIA)

Contact me if you need this

Parallel Computing

Using multiple processing units in parallel to solve problems more quickly than with a single processing unit

Applications in engineering and design (DNA sequence analysis)

Scientific applications (oceanography, astrophysics)

Commercial applications (data mining, transaction processing)

Why More Speed or Parallelism

Scientific applications

Computational model to simulate the impact of climate change

Video and audio coding and manipulation

UHD TV

User interfaces

Modern smart phone users enjoy a more natural interface with high-resolution touch screens

Big data applications

Processing huge amount of data

Single Processor Performance

From 1986-2002, microprocessors' performance was increasing an average of 50% per year

Then single-processor performance dropped to about 20% increase per year

Increase in single processor performance has been driven by the increasing *density*, the decreasing *size* of transistors

Moore's Law


An observation of Gordon Moore in 1965, that the number of transistors in a dense integrated circuit doubles approximately every two years

Moore's Law: The number of transistors on microchips doubles every two years

Moore's law describes the empirical regularity that the number of transistors on integrated circuits doubles approximately every two years. This advancement is important for other aspects of technological progress in computing – such as processing speed or the price of computers.

Our World
in Data

Transistor count



Data source: Wikipedia (wikipedia.org/wiki/Transistor_count)

OurWorldInData.org – Research and data to make progress against the world's largest problems.

Licensed under CC-BY by the authors Hannah Ritchie and Max Roser.

Cooking-Aware Computing



Power Wall

More/smaller transistors = faster processors

Faster processors = increased power consumption

Increased power consumption = increased heat

Increased heat = unreliable processors

Solution: Move away from single-core systems to multiprocessor systems

Multicore Systems

CPU: The "Central Processing Unit"


Traditionally, applications use CPU for primary calculations

General-purpose capabilities

Established technology


Usually equipped with 8 or less powerful cores

Optimal for concurrent processes but not large scale parallel computations



Multicore Systems

Instead of designing and building faster microprocessors, put multiple processors (each core is less powerful than the previous generation's single core design) on a single integrated circuit



Manycore Systems

GPU: The "Graphics Processing Unit"

Relatively new technology designed for parallelizable problems

Initially created specifically for graphics

Became more capable of general computations

GeForce 256, world's first GPU, in 1999 by NVIDIA



GPUs


Everywhere from inside smartphones, laptops, datacenters, and all the way to supercomputers

Initially, real-time rendering with a focus on video games


Increasingly support non-graphics computing

Refined GPU architectures and programming model to increase flexibility as well as energy efficiency

Heterogeneous Parallel Computing





CPU vs GPU Performance



Graphics

Example: Raytracing

Enables real-time light reflections and cinematic effects in games



<https://www.youtube.com/watch?v=gq1klB4inoc>

General-Purpose Computation on GPUs

The GPU is no longer just for graphics

Particle systems, collision detection


Fluid dynamics

Simulation

Neural network

High Performance Computing

ACCELERATING HPC



All results are measured
Except BerkeleyGW, V100 used is single V100 SXM2. A100 used is single A100 SXM4
More apps detail: AMBER based on PME-Cellulose, GROMACS with STMV (h-bond), LAMMPS with Atomic Fluid LJ-2.5, NAMD with v3.0a1 STMV_NVE
Chroma with szsc121_24_128, FUN3D with dpw, RTM with Isotropic Radius 4 1024³, SPECFEM3D with Cartesian four material model
BerkeleyGW based on Chi Sum and uses 8xV100 in DGX-1, vs 8xA100 in DGX A100

Deep Learning


Why neural networks become very popular in recent years, though much theory was developed 20 years ago?

Huge amount of data

Computational power

<https://www.youtube.com/watch?v=DiGB5uAYKAg>


Deep Learning



Deep Learning

NVIDIA AI and H100 Deliver 6.7X in 2.5 Years

Full-stack innovation fuels continuous performance gains



ResNet-50 v1.5: 8x NVIDIA 0.7-18, 8x NVIDIA 2.1-2060, 8x NVIDIA 2.1-2091 | BERT: 8x NVIDIA 0.7-19, 8x NVIDIA 2.1-2062, 8x NVIDIA 2.1-2091 | DLRM: 8x NVIDIA 0.7-17, 8x NVIDIA 2.1-2059, 8x NVIDIA 2.1-2091 | Mask R-CNN: 8x NVIDIA 0.7-19, 8x NVIDIA 2.1-2062, 8x NVIDIA 2.1-2091 | RetinaNet: 8x NVIDIA 2.0-2091, 8x NVIDIA 2.1-2061, 8x NVIDIA 2.1-2091 | RNN-T: 8x NVIDIA 1.0-1060, 8x NVIDIA 2.1-2061, 8x NVIDIA 2.1-2091 | Mini Go: 8x NVIDIA 0.7-20, 8x NVIDIA 2.1-2063, 8x NVIDIA 2.1-2091 | 3D U-Net: 8x NVIDIA 1.0-1059, 8x NVIDIA 2.1-2060, 8x NVIDIA 2.1-2091
First NVIDIA A100 Tensor Core GPU results normalized for throughput due to higher accuracy requirements introduced in MLPerf™ Training 2.0 where applicable.
MLPerf™ name and logo are trademarks. See www.mlperf.org for more information.

Bitcoin Mining

HOW DO BITCOINS WORK?



'Miners' create Bitcoins by using computers to solve mathematical functions. The same process also verifies previous transactions



Bitcoin exchanges will trade between conventional currencies and Bitcoin, offering a way into the market for non-miners, as well as a way to cash out



Users download a Bitcoin 'wallet' that works a little like an email address, providing a way to store and receive currency. Bitcoins can be transferred from one wallet to another using a web browser or

Businesses create a wallet in the same way as an individual user, typically using a website button to enable a Bitcoin payment. For in-the-flesh enterprises, QR codes can be used to let customers pay quickly



Embedded Systems



Apple A17 Bionic

Mobile chip inside iPhone 15

**64-bit ARM-based system on a chip (SoC)
designed by Apple**

6-core multicore

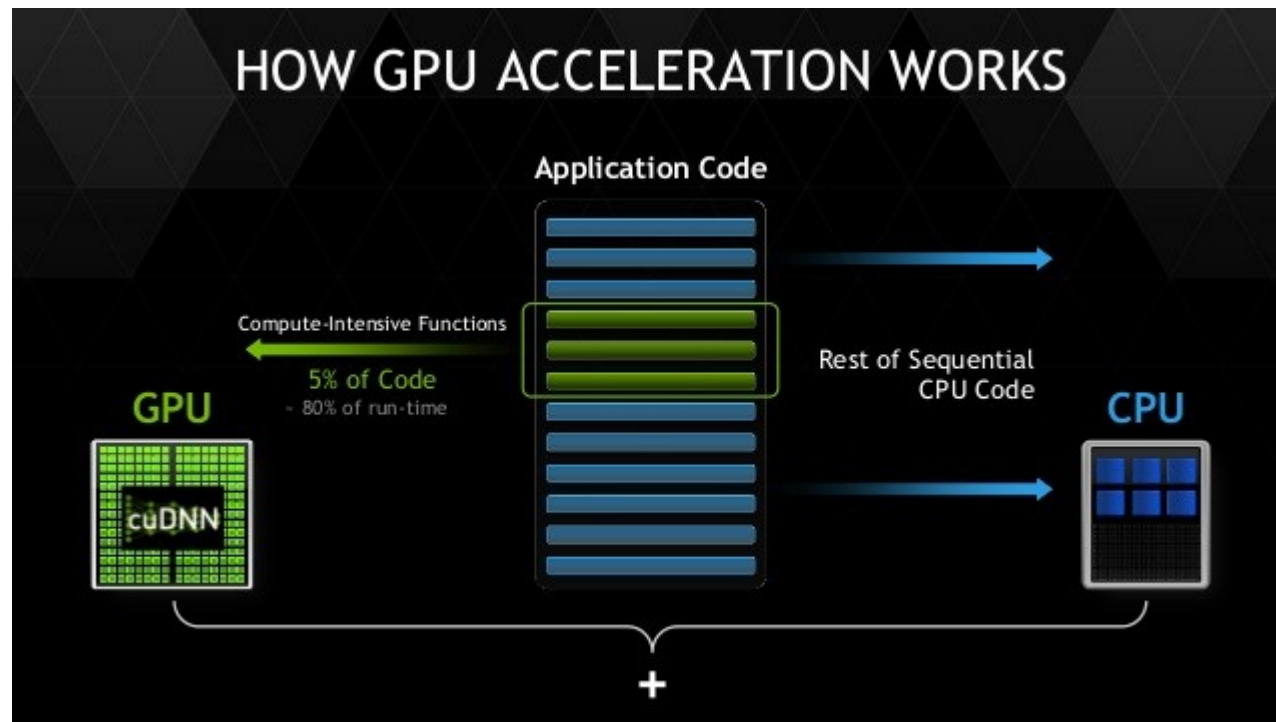
Two high-performance cores (3.78 GHz)

Four low-power energy-efficient cores (2.11 GHz)

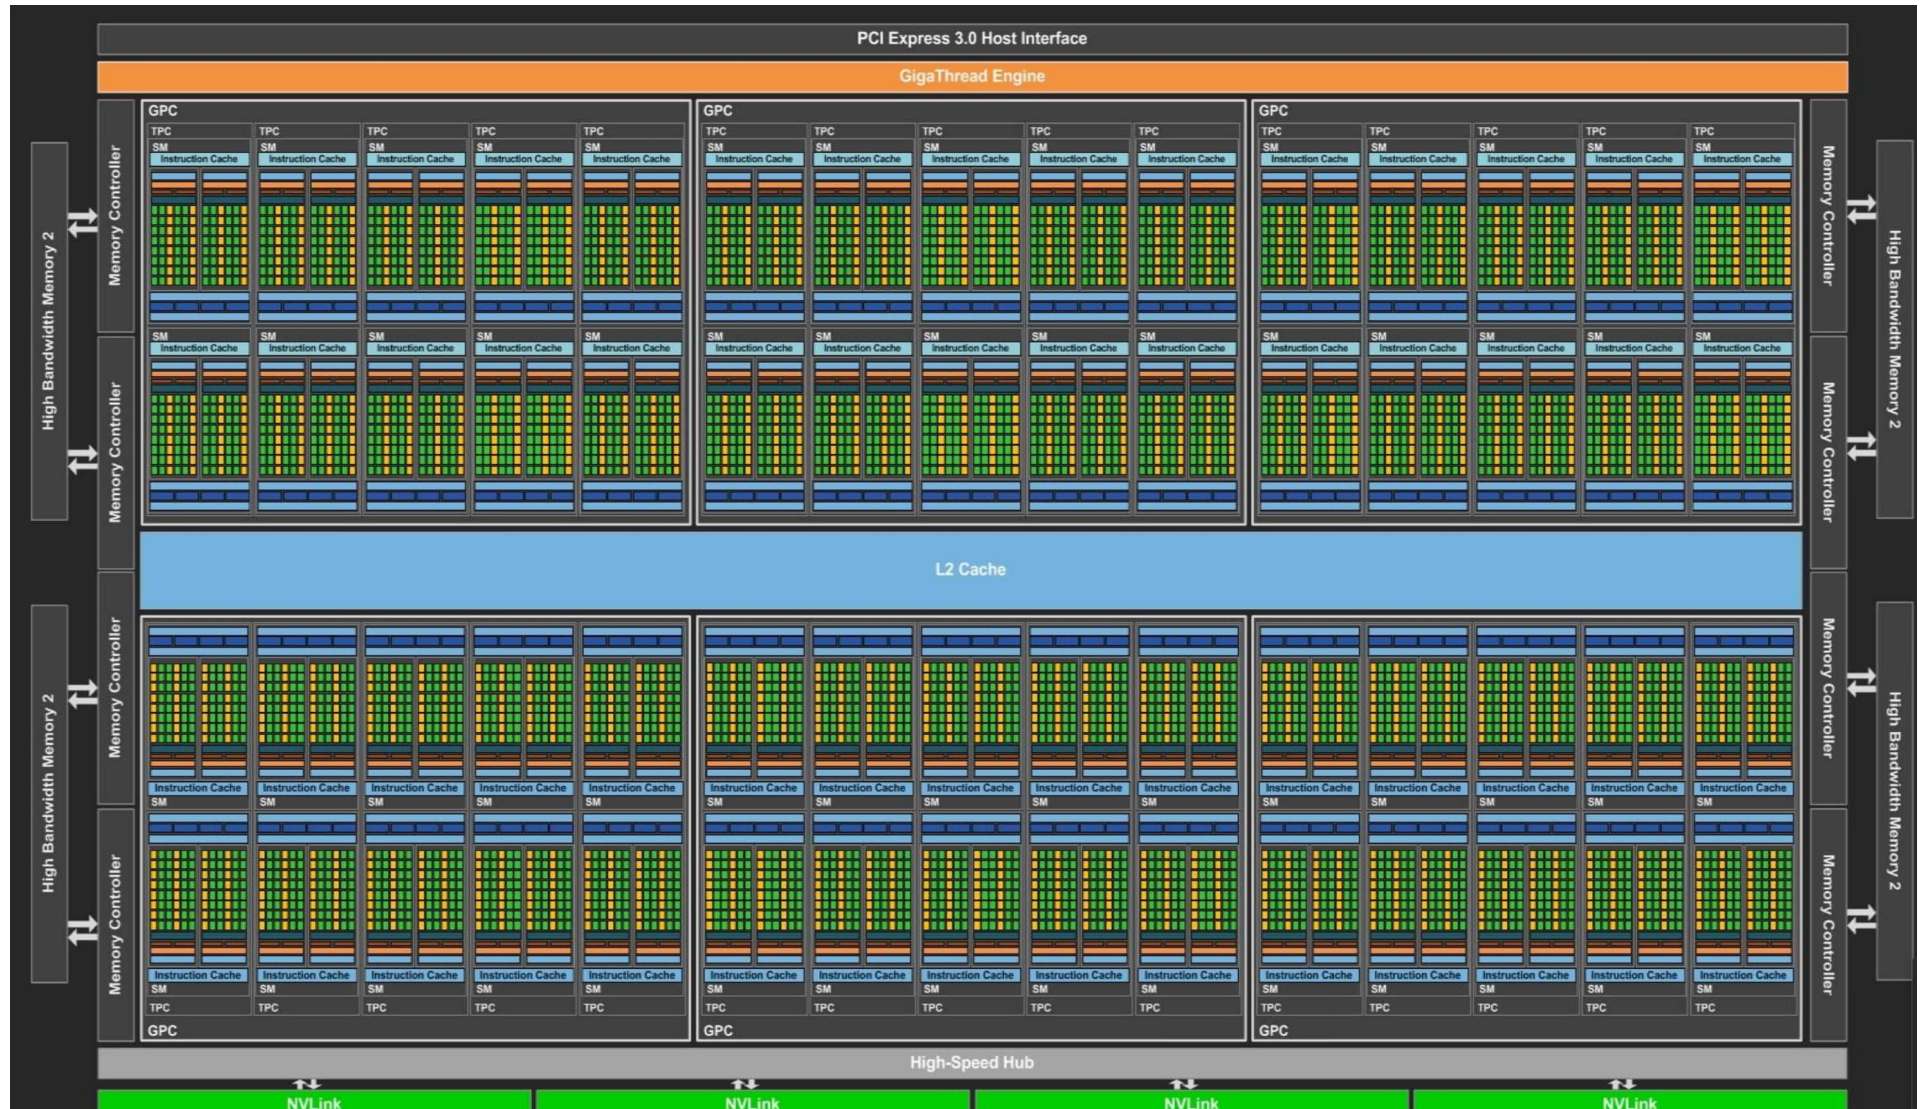
6-core GPU

16-core neural engine

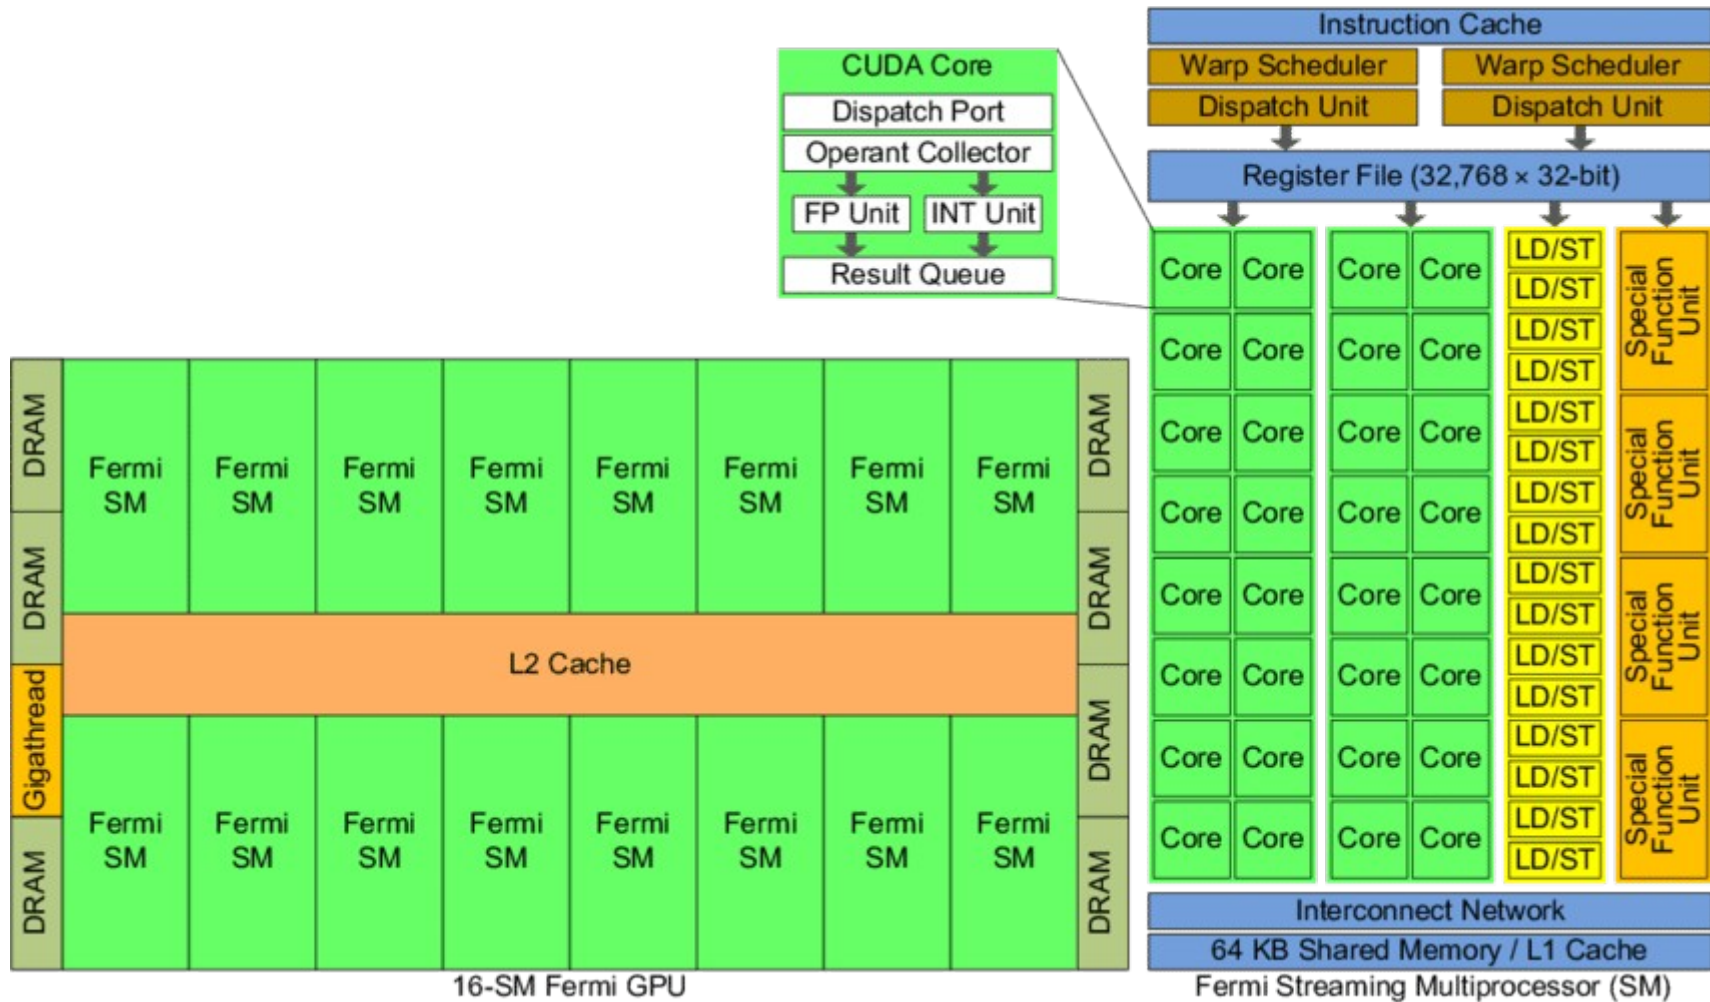
GPU Acceleration



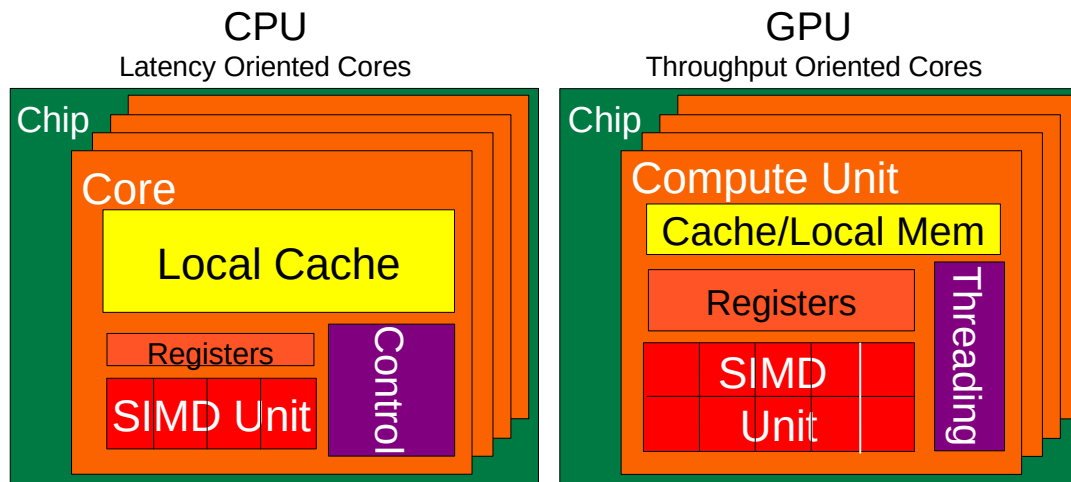
GPU Architecture



GPU Architecture Example



CPU vs GPU



Latency vs Throughput

Latency

Elapsed time

Throughput

Events per unit time


Example: perform 100 multiplications

Single-core CPU: 1 mul = 1 sec, 100 mul = 100 sec

Quad-core CPU: 1 mul = 2 sec, 100 mul = 50 sec

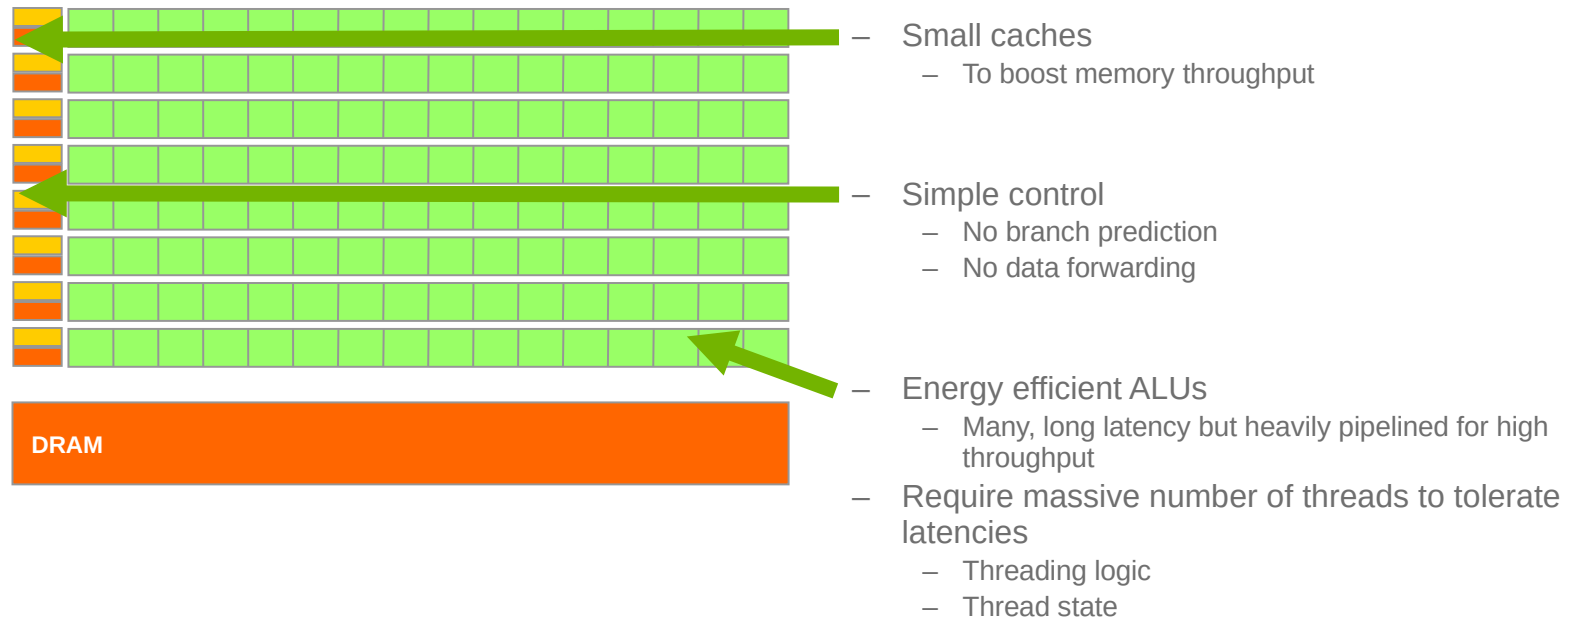
GPU (50 PUs): 1 mul = 5 sec, 100 mul = 10 sec

CPUs: Latency Oriented Design



minimize the execution latency of a single thread

GPUs: Throughput Oriented Design



Winning Applications with Both CPU and GPU

CPUs for sequential parts where latency matters

CPUs can be 10X+ faster than GPUs for sequential code

GPUs for parallel parts where throughput wins

GPUs can be 10X+ faster than CPUs for parallel code

Problems with GPUs

Need enough parallelism

Data-centric applications

Under-utilization

Memory bottleneck

Bandwidth to CPU

Separate memories

GPU Advantages

a very large presence in the market place

Only a few elite applications funded by government and large corporations have been successfully developed on traditional parallel computing systems

GPUs have been sold by the hundreds of millions

practical form factors and easy accessibility

for medical imaging, fine to publish a paper based on a 64-node cluster machine

But real-world clinical applications on MRI machines utilize some combination of a PC and special hardware accelerators

GPU Advantages

executing numeric computing applications

IEEE Floating-Point Standard was not strong in early GPUs

Up to 2009, a major barrier was that the GPU floating-point arithmetic units were primarily single precision

this has changed with the recent GPUs whose double precision execution speed approaches about half that of single precision, a level that only high-end CPU cores achieve

programming effort

Until 2006, OpenGL or Direct3D techniques were needed

Much easier with the release of CUDA

Simple Example

Add two arrays

$A[\] + B[\] \rightarrow C[\]$

CPU:

```
float *C = malloc(N * sizeof(float));
```

```
for (int i = 0; i < N; i++)
```

```
    C[i] = A[i] + B[i];
```

```
return C;
```

Parallel Version

(allocate memory for C)

Create # of threads equal to number of cores on processor (around 2, 4, perhaps 8)

(Indicate portions of A, B, C to each thread...)

In each thread,

For (i from beginning region of thread)

$C[i] \leftarrow A[i] + B[i]$

//lots of waiting involved for memory reads, writes, ...

Wait for threads to synchronize...

This is slightly faster - 2-8x (slightly more with other tricks)

Parallel Performance

How many threads? How does performance scale?

Context switching:

The action of switching which thread is being processed

High penalty on the CPU

Not an issue on the GPU

GPU Version

(allocate memory for A, B, C on GPU)

Create the “kernel” – each thread will perform one (or a few) additions

Specify the following kernel operation:

For all i's (indices) assigned to this thread:

$C[i] \leftarrow A[i] + B[i]$

Start ~20000 (!) threads

Wait for threads to synchronize...

GPU Performance

We have lots of cores

This allows us to run many threads simultaneously with no context switches

In a typical system, thousands of threads are queued up for work. If the GPU must wait on one group of threads, it simply begins executing work on another.

Parallel Programming

Design parallel algorithms with the same level of algorithmic (computational) complexity as sequential algorithms (large parallel overheads)

The execution speed of many applications is limited by memory access speed

The execution speed of parallel programs is often more sensitive to the input data characteristics

Parallel Programming Models

Distributed-memory programming

MPI, PVM

Shared-memory programming

Pthreads, OpenMP

GPU programming

CUDA, OpenCL

MapReduce programming

Hadoop

CUDA

CUDA (Compute Unified Device Architecture)

Enables a general purpose programming model on NVIDIA GPUs

Enables explicit GPU memory management

GPU threads are extremely lightweight

Very little creation overhead

GPU needs 1000s of threads for full efficiency

Multi-core CPU needs only a few

CUDA Program

CUDA platform is accessible through CUDA-accelerated libraries and APIs

CUDA C is an extension of standard ANSI C with language extensions to enable heterogeneous programming, and also APIs to manage devices, memory, and other tasks

Integrated host+device app C program

Serial or modestly parallel parts in host C code

Highly parallel parts in device SPMD kernel C code